

**ACCELERATOR BASED PROGRAMMING**  
**UPPSALA UNIVERSITY**  
**FALL 2023**

EXERCISE 4: FIRST STEPS IN KOKKOS

This exercise is a preparation for the fourth assignment.

**1. Exercise Goal.** The goal of this exercise is to get started with Kokkos. The goal is to run the exercises 01 and 02 listed at <https://github.com/kokkos/kokkos-tutorials/tree/main/Exercises> on both the CPU and the GPU of UPPMAX/Snowy.

The tasks are as follows:

- Familiarize yourself with Kokkos by looking at the lecture material, and in particular the Kokkos guide [https://github.com/kokkos/kokkos-tutorials/blob/main/LectureSeries/KokkosTutorial\\_01\\_Introduction.pdf](https://github.com/kokkos/kokkos-tutorials/blob/main/LectureSeries/KokkosTutorial_01_Introduction.pdf).
- On your computer or UPPMAX, go to your favorite directory, e.g.  
  `${HOME}/Kokkos`  
and download Kokkos:  
  `git clone https://github.com/kokkos/kokkos`
- Download the Kokkos tutorials  
  `git clone https://github.com/kokkos/kokkos-tutorials`
- Download the Kokkos tutorials
- We also need to load the compilers for gcc and nvidia  
  `module load gcc/8.4.0`  
  `module use /sw/EasyBuild/snowy/modules/all/`  
  `module use /sw/EasyBuild/snowy-GPU/modules/all/`  
  `module use python_ML_packages/3.9.5-gpu`  
  `module load CUDA/11.7.0`
- Go to the first exercise  
  `cd kokkos-tutorials/Exercises/01/Begin`
- Inspect the `Makefile` and adjust the paths to your system. If you have Kokkos in a directory `Kokkos/kokkos` of your home directory, you do not have to adjust it.
- To compile for the CPU, you do not have to do further adjustments and you can just run `make -j8`. To compile for the GPU, you need to either pass options to the command line when calling `make` or adjust the settings in the `Makefile`. To compile for CUDA, change the second line to

```
KOKKOS_DEVICES = "Cuda"
```

and line 14 to

```
KOKKOS_ARCH = "TURING75"
```

Also, to code on the CPU nodes of Snowy you need to adjust the CPU architecture as well, line 19,

```
KOKKOS_ARCH = "SNB"
```

- Follow *exercise 1* and initialize Kokkos. Compare with the respective code in the solution. This is a CPU only code we parallelize with OpenMP. A possible runscript could look like this

```
#!/bin/bash -l

#SBATCH -A uppmx2023-2-36
#SBATCH -M snowy
#SBATCH -p node
#SBATCH -N 1
#SBATCH -t 0:10:00
#SBATCH -J Exercise1
#SBATCH -D ./

export OMP_PROC_BIND=spread OMP_PLACES=threads
echo "8 threads"
export OMP_NUM_THREADS=8
./01_Exercise.host -N 13
./01_Exercise.host -N 14
./01_Exercise.host -N 15
./01_Exercise.host -N 16

echo "16 threads"
export OMP_NUM_THREADS=16

./01_Exercise.host -N 13
```

- Work on *exercise 2* regarding the memory spaces. Note that this will use unified memory between CPU and GPU, as seen by the `force_uvm` CUDA option. Run the CUDA code for this example and compare bandwidth for various sizes of the matrix